



Anisha
Arkadip Basu
Naveen Rupdrapa
Yatish Reddy

Q-Learning with UCB Exploration

IISc, M.Tech (Online) Aug 2022, EA-277o

11 Dec 2022

Motivation

- To improve the performance of **Model Free RL algorithm** for infinite horizon MDP
- Deal with infinite horizon MDPs without access to simulator
- Before the paper "Q-LEARNING WITH UCB EXPLORATION IS SAMPLE EFFICIENT FOR INFINITE-HORIZON MDP" was published. The best known sample complexity of exploration was achieved by delayed Q-learning
$$\tilde{O}\left(\frac{SA}{\epsilon^4(1-\gamma)^8}\right)$$

Big O Tilde, log factors can be ignored
- *Several model-based algorithms have been proposed for infinite horizon MDP. However, there still exists a considerable gap between the state-of-the-art algorithm and the theoretical lower bound regarding $1/(1-\gamma)$ factor*
- *The performance measure cannot be a straightforward extension of the sample complexity defined by finite horizon setting.*

Q-Learning

Q-learning lets the agent use the environment's rewards to learn, over time, the best action to take in a given state.

The values store in the Q-table are called a Q-values, and they map to a (state, action) combination.

Q-values are initialized to an arbitrary value, and as the agent exposes itself to the environment and receives different rewards by executing different actions, the Q-values are updated using the equation:

$$Q(state, action) \leftarrow (1 - \alpha)Q(state, action) + \alpha \left(reward + \gamma \max_a Q(next\ state, all\ actions) \right)$$

α (alpha) is the learning rate, γ (gamma) is the discount factor.

After enough random exploration of actions, the Q-values tend to converge serving our agent as an action-value function which can be exploited to pick the most optimal action from a given state.

Q learning with UCB exploration

Q-learning Algorithm:

- ▶ For $t = 1, 2, \dots$
 - ▶ Act $a_t = \arg \max_a Q(s_t, a)$,
 - ▶ $k \leftarrow$ number of times (s_t, a_t) is visited,
 - ▶ $Q(s_t, a_t) \leftarrow (1 - \alpha_k)Q(s_t, a_t) + \alpha_k [r_t + \gamma V(s_{t+1})]$.

Q-learning with Hoeffding-style UCB exploration bonus:

- ▶ For $t = 1, 2, \dots$
 - ▶ Act $a_t = \arg \max_a Q(s_t, a)$,
 - ▶ $Q(s_t, a_t) \leftarrow (1 - \alpha_k)Q(s_t, a_t) + \alpha_k \left[r_t + \gamma V(s_{t+1}) + \tilde{\Theta} \left(\sqrt{\frac{1}{(1-\gamma)^3 k}} \right) \right]$.

Terminologies

- **Sample complexity of Exploration:** Sample complexity of Exploration of an algorithm is defined as the number of time steps t such that the non-stationary policy π_t at time t , is not ϵ -optimal for current state s_t .

$$V^{\pi_t}(s_t) < V^*(s_t) - \epsilon.$$

- **Probably Approximately Correct in Markov Decision Processes (PAC-MDP):** An algorithm is said to be PAC-MDP if, for any ϵ and δ , the sample complexity of ALG is less than some polynomial in the relevant quantities $(S, A, 1/\epsilon, 1/\delta, 1/(1 - \gamma))$ with probability at least $1 - \delta$.
- **Bellman equation:** long-term- reward in a given action is equal to the reward from the current action combined with the expected reward from the future actions taken at the following time.

$$\begin{cases} V^{\pi_t}(s) = Q^{\pi_t}(s, \pi_t(s)) \\ Q^{\pi_t}(s, a) := (r_t + \gamma \mathbb{P} V^{\pi_{t+1}})(s, a), \end{cases} \quad \begin{cases} V^*(s) = Q^*(s, \pi^*(s)) \\ Q^*(s, a) := (r_t + \gamma \mathbb{P} V^*)(s, a), \end{cases}$$

Complexity measurement

- The main bottleneck in the infinite horizon setting, the agent may enter under-explored regions at any time period, and sample complexity of exploration characterizes the performance at all states the agent enters.
- First we need to establish convenient sufficient conditions for being ϵ -optimal at timestep t and state s_t
i.e.

$$V^*(s_t) - V^{\pi_t}(s_t) \leq \epsilon.$$

Infinite Q-learning with UCB

3.1 ALGORITHM

Algorithm 1 Infinite Q-learning with UCB

Parameters: ϵ, γ, δ

Initialize $Q(s, a), \hat{Q}(s, a) \leftarrow \frac{1}{1-\gamma}, N(s, a) \leftarrow 0, \epsilon_1 \leftarrow \frac{\epsilon}{24RM \ln \frac{1}{1-\gamma}}, H \leftarrow \frac{\ln 1/((1-\gamma)\epsilon_1)}{\ln 1/\gamma}$.

Define $\iota(k) = \ln(SA(k+1)(k+2)/\delta), \alpha_k = \frac{H+1}{H+k}$.

for $t = 1, 2, \dots$ **do**

5: Take action $a_t \leftarrow \arg \max_{a'} \hat{Q}(s_t, a')$

 Receive reward r_t and transit to s_{t+1}

$N(s_t, a_t) \leftarrow N(s_t, a_t) + 1$

$k \leftarrow N(s_t, a_t), b_k \leftarrow \frac{c_2}{1-\gamma} \sqrt{\frac{H\iota(k)}{k}} \quad \triangleright c_2 \text{ is a constant and can be set to } 4\sqrt{2}$

$\hat{V}(s_{t+1}) \leftarrow \max_{a \in A} \hat{Q}(s_{t+1}, a)$

10: $Q(s_t, a_t) \leftarrow (1 - \alpha_k)Q(s_t, a_t) + \alpha_k [r(s_t, a_t) + b_k + \gamma \hat{V}(s_{t+1})]$

$\hat{Q}(s_t, a_t) \leftarrow \min(\hat{Q}(s_t, a_t), Q(s_t, a_t))$

end for

UCB Q-learning algorithm (Algorithm 1) maintains an optimistic estimation of action value function $Q(s, a)$ and its historical minimum value $\hat{Q}(s, a)$.

Complexity of the Algorithm

- With some mathematical theorems and assumptions it has been proved in the paper that with probability $1 - \delta$, the number of time steps such that $(V^* - V^\pi)(s_t) > \epsilon$ is


$$\tilde{O} \left(\frac{SA \ln 1/\delta}{\epsilon^2 (1 - \gamma)^7} \right)$$

- The complexity is better than delayed q learning.

Demo : Taxi problem

There are 500 discrete states since there are 25 taxi positions, 5 possible locations of the passenger (including the case when the passenger is the taxi), and 4 destination locations

Rewards: There is a reward of -1 for each action and an additional reward of +20 for delivering the passenger. There is a reward of -10 for executing actions "pickup" and "dropoff" illegally.
Rendering:



- The filled square represents the taxi, which is yellow without a passenger and green with a passenger.
- The pipe ("|") represents a wall which the taxi cannot cross.
- R, G, Y, B are the possible pickup and destination locations.
- Action Space: 6, State Size: 500

Actions:

- 0: move south
- 1: move north
- 2: move east
- 3: move west
- 4: pickup passenger
- 5: dropoff passenger

Fine tuning

- $\frac{1}{2} < \gamma < 1$
- $R = (\log(3 / (\epsilon_{\text{hyp}} * (1 - \gamma)))) / (1 - \gamma)$
- $M = 2 * \log_2(1 / ((1 - \gamma) * \epsilon_{\text{hyp}}))$
- $\epsilon_1 = (\epsilon_{\text{hyp}} / (24 * R * M * \log(1 / (1 - \gamma))))$
- $H = (\log(1 / (\epsilon_1 * (1 - \gamma)))) / \log(1 / \gamma)$
- `qcap_table = np.full((state_size, action_size), q_initial_value)`
- `n_table = np.zeros((state_size, action_size))`
- `q_table = np.zeros((state_size, action_size))`
- `max_epsilon = 1.0`
- `min_epsilon = 0.01`
- `decay_rate = 0.01`
- `gamma = 0.9`
- `c2 = 4 * math.sqrt(2)`

Changes made to Algorithm

```
for episode in range(total_ep):
    state = env.reset()
    step = 0
    done = False

    for step in range(max_steps):

        exp_exp_tradeoff = random.uniform(0,1)
        #Take action at  $\leftarrow \arg \max_a Q^*(st, a)$ 
        if exp_exp_tradeoff > epsilon:
            action = np.argmax(qcap_table[state, :])
        else:
            action = env.action_space.sample()

        # Receive reward rt and transit to st+1
        new_state, reward, done, info = env.step(action)

        #N(st, at)  $\leftarrow N(st, at) + 1$ 
        n_table[state, action] = n_table[state, action] + 1

        #k  $\leftarrow N(st, at)$ 
        k = n_table[state, action]

        #update bk
        bk = (c2 / (1 - gamma)) * math.sqrt((H * getlk(k, delta)) / k)

        #action_max = np.argmax(qcap_table[new_state, :])
        vcapstplus1 = np.max(qcap_table[new_state, :])
        # Update  $Q(s,a) := Q(s,a) + lr [R(s,a) + gamma * \max_{a'} Q(s',a') - Q(s,a)]$ 
        alphak = getaplhak(H, k)
        q_table[state, action] = (1 - alphak) * q_table[state, action] + alphak * (reward + bk + gamma * vcapstplus1)

        # $Q^*(st, at) \leftarrow \min(Q^*(st, at), Q(st, at))$ 

        qcap_table[state, action] = min(q_table[state, action], qcap_table[state, action])
        state = new_state
        if done == True:
            break

    episode += 1

    epsilon = min_epsilon + (max_epsilon - min_epsilon) * np.exp(-decay_rate * episode)
```

Colab results

Total Episodes	Mean Score
15000	-1093.4
30000	-858.5
50000	-712.2
200000	-522.6

Reference

1. <https://arxiv.org/abs/1901.09311>
2. <https://openreview.net/pdf?id=BkgISTNFDB>
4. <https://paperswithcode.com/task/q-learning>
5. <https://blog.paperspace.com/getting-started-with-openai-gym/>
6. [https://www.researchgate.net/publication/335805245_Q-Learning Algorithms A Comprehensive Classification and Applications](https://www.researchgate.net/publication/335805245_Q-Learning_Algorithms_A_Comprehensive_Classification_and_Applications)

Code

<https://www.kaggle.com/anishabhushan/q-learning-taxi-implementation-ucb-final>

Anisha
Arkadip Basu
Naveen Rupdrapa
Yatish Reddy

IISc, M.Tech (Online)
Reinforcement Learning
EA-277o Aug 2022 batch



Thank You!